

# MATH 3190 Homework 6

Focus: Notes 8

Due March 30, 2024

Your homework should be completed in R Markdown or Quarto and Knitted to an html or pdf document. You will “turn in” this homework by uploading to your GitHub Math\_3190\_Assignment repository in the Homework directory.

Some of the parts in problems 1 and 2 require writing down some math-heavy expressions. You may either type it up using LaTeX style formatting in R Markdown, or you can write it by hand (neatly) and include pictures or scans of your work in your R Markdown document.

## Problem 1 (10 points)

Three airlines serve a small town in Ohio. Airline A has 52% of all scheduled flights, airline B has 35% and airline C has the remaining 13%. Their on-time rates are 85%, 67%, and 41%, respectively. A flight just left on-time. What is the probability that it was a flight of airline A?

**The method I used here is based on the example from page 4 of Notes 8.**

$$\begin{aligned} & \frac{(0.52)(0.85)}{(0.35)(0.67) + (0.52)(0.85) + (0.13)(0.41)} \\ &= \frac{0.442}{0.7298} \\ &= 0.6056 \end{aligned}$$

## Problem 2 (13 points)

Suppose we have a data set with each observation  $x_i$  independent and identically exponentially distributed for  $i = 1, 2, \dots, n$ . That is,  $x_i \sim \text{Exp}(\lambda)$  where  $\lambda$  is the rate parameter. We would like to find a posterior (or at least a function proportional to it) for  $\lambda$ .

### Part a (5 points)

Write down the likelihood function (or a function proportional to it) in this situation. We would call this  $p(x|\lambda)$ .

$$p(x|\lambda) = \prod_{i=1}^n \lambda e^{-\lambda x_i}$$

### Part b (5 points)

Now let  $\lambda$  have a normal prior with mean 0.1 and variance 1:  $\lambda \sim N(1/10, 1)$ . Use this and the likelihood from part a to write down a function that is proportional to the posterior of  $\lambda$  given  $\mathbf{x}$ . We call this  $p(\lambda|\mathbf{x})$ .

$$\exp\left\{-\frac{1}{2}(\lambda - 0.1)^2\right\} * \prod_{i=1}^n \lambda e^{-\lambda x_i}$$

### Part c (3 points)

Which would be more appropriate here to obtain samples of  $\lambda$ , the Gibbs or Metropolis algorithm? Explain why. You may want to look on page 8 of Notes 8 in the conjugate prior table.

Metropolis sampling, because we do not have a conjugate prior

## Problem 3 (26 points)

Suppose we have the vector  $\mathbf{x} = c(1.83, 1.72, 2.13, 2.49, 0.90, 2.01, 1.51, 3.12, 1.29, 1.54, 2.94, 3.02, 0.93, 2.78)$  that we believe comes from a gamma distribution with shape of 10 and some rate  $\beta$ :  $x_i \sim \text{Gam}(10, \beta)$ . We will use sampling to obtain some information about  $\beta$ . Let's put a gamma prior on  $\beta$  with a shape of  $\alpha_0$  and a rate of 1:  $\beta \sim \text{Gam}(\alpha_0, 1)$ .

### Part a (5 points)

Use the fact that this is a conjugate prior to write down what kind of distribution the posterior of  $\beta$ , which is  $p(\beta|\mathbf{x})$ , is.

```
x <- c(1.83, 1.72, 2.13, 2.49, 0.90, 2.01, 1.51, 3.12, 1.29, 1.54, 2.94, 3.02, 0.93, 2.78)
n <- length(x)
beta0 <- 1
alpha <- 10
rate <- beta0 + sum(x)
shape <- n*alpha
```

$$p(\beta|\mathbf{x}) \sim \text{gam}(\alpha_0 + 140, 29.21)$$

### Part b (5 points)

Let  $\alpha_0 = 1$ . In an **R** code chunk, sample 10,000  $\beta$  values from the distribution you wrote down in part a using the `rgamma()` function and report the 95% credible interval for  $\beta$  using the 2.5th and 97.5th percentiles.

```
library(tidyverse)
```

```
## Warning: package 'tidyverse' was built under R version 4.3.2
```

```
## Warning: package 'ggplot2' was built under R version 4.3.2
```

```
## Warning: package 'tibble' was built under R version 4.2.3
```

```
## Warning: package 'tidyr' was built under R version 4.3.2
```

```
## Warning: package 'readr' was built under R version 4.3.2
```

```
## Warning: package 'purrr' was built under R version 4.3.2
```

```
## Warning: package 'dplyr' was built under R version 4.2.3
```

```
## Warning: package 'stringr' was built under R version 4.3.2

## Warning: package 'forcats' was built under R version 4.3.2

## Warning: package 'lubridate' was built under R version 4.3.2

## -- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
## v dplyr      1.1.4      v readr      2.1.5
## v forcats    1.0.0      v stringr   1.5.1
## v ggplot2    3.4.4      v tibble    3.2.1
## v lubridate  1.9.3      v tidyr     1.3.1
## v purrr      1.0.2
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
```

```
credInts1 <- tibble(alpha0 = numeric(),
                    lBound = numeric(),
                    uBound = numeric())

alphas <- c(1, 10, 100)
alpha <- 10
beta <- 1
for (i in alphas){
  shape <- i + length(x)*alpha
  rate <- beta0 + sum(x)
  beta <- rgamma(10000, shape, rate)
  credInt <- quantile(beta, c(0.025, 0.975))
  credInts1 <- add_row(credInts1,
                      alpha0 = i,
                      lBound = credInt[1],
                      uBound = credInt[2])
}
print(credInts1)
```

```
## # A tibble: 3 x 3
##   alpha0 lBound uBound
##   <dbl> <dbl> <dbl>
## 1      1  4.06  5.66
## 2     10  4.34  5.99
## 3    100  7.21  9.27
```

### Part c (3 points)

Repeat part b with  $\alpha_0 = 10$ .

See credInts1 table from part a

### Part d (3 points)

Repeat part b with  $\alpha_0 = 100$ .

See `credInts1` table from part a

### Part e (7 points)

Now suppose we have twice as much data (given in the **R** code chunk below). Repeat parts b, c, and d using this `x` vector instead and report the three 95% credible intervals. Note, this new vector `x` will change the shape and rate parameters used in the `rgamma()` functions.

```
x <- c(1.83, 1.72, 2.13, 2.49, 0.90, 2.01, 1.51, 3.12, 1.29, 1.54,
      2.94, 3.02, 0.93, 2.78, 2.76, 1.70, 1.42, 2.16, 1.07, 2.21,
      2.38, 2.27, 1.72, 1.44, 1.54, 1.72, 1.87, 1.39)
```

```
credInts2 <- tibble(alpha0 = numeric(),
                    lBound = numeric(),
                    uBound = numeric())

alphas <- c(1, 10, 100)
alpha <- 10
beta <- 1
for (i in alphas){
  shape <- i + length(x)*alpha
  rate <- beta0 + sum(x)
  beta <- rgamma(10000, shape, rate)
  credInt <- quantile(beta, c(0.025, 0.975))
  credInts2 <- add_row(credInts2,
                      alpha0 = i,
                      lBound = credInt[1],
                      uBound = credInt[2])
}
print(credInts2)
```

```
## # A tibble: 3 x 3
##   alpha0 lBound uBound
##   <dbl> <dbl> <dbl>
## 1      1  4.54  5.73
## 2     10  4.70  5.91
## 3    100  6.24  7.64
```

### Part f (3 points)

In this problem, the true  $\beta$  value is 5. Write a sentence or two about the effect adding more data has to these credible intervals by comparing the intervals from parts b-d to the intervals from part e.

**It looks like adding more data made the bounds of the credible intervals narrower. The intervals made with  $\alpha = 1$  and  $\alpha = 10$  both contain 5, but the ones with  $\alpha = 100$  do not contain 5. It seems like adding more data makes the intervals more precise for all values of  $\alpha$ .**

## Problem 4 (51 points)

Let's apply the Bayesian framework to a regression problem. In the GitHub data folder, there is a file called `treeseeds.txt` that contains information about species of tree, the count of seeds it produces, and the average weight of those seeds in mg.

### Part a (3 points)

Read in the `treeseeds.txt` file and take the log of the counts and weights. Fit an OLS regression model using  $\log(\text{weight})$  to predict  $\log(\text{count})$ .

```
tree <- read_csv("treeseeds.txt")

## Rows: 19 Columns: 3
## -- Column specification -----
## Delimiter: ","
## chr (1): species
## dbl (2): count, weight
##
## i Use 'spec()' to retrieve the full column specification for this data.
## i Specify the column types or set 'show_col_types = FALSE' to quiet this message.
```

```
tree['count'] = log(tree['count'])
tree['weight'] = log(tree['weight'])
```

```
treeMod <- lm(count ~ weight, data = tree)
summary(treeMod)
```

```
##
## Call:
## lm(formula = count ~ weight, data = tree)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.4429 -0.3877  0.1778  0.6167  1.2691
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   9.37903    0.39720  23.613 1.95e-14 ***
## weight       -0.51491    0.07185  -7.166 1.58e-06 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.9363 on 17 degrees of freedom
## Multiple R-squared:  0.7513, Adjusted R-squared:  0.7367
## F-statistic: 51.35 on 1 and 17 DF,  p-value: 1.58e-06
```

### Part b (15 points)

We will walk through the mathematics of obtaining the posterior together here since this problem will focus on coding the Metropolis algorithm. Assuming the true errors are normal with mean 0 and variance  $\sigma^2$ ,

$\epsilon_i \sim N(0, \sigma^2)$ , it can be shown that each  $y_i$  has the distribution

$$p(y_i|x_i, \beta_0, \beta_1 \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{1}{2\sigma^2}(y_i - \beta_0 - \beta_1 x_i)^2\right)$$

So, we can write the likelihood is

$$p(y_i|x_i, \beta_0, \beta_1 \sigma^2) \propto \exp\left(-\frac{1}{2\sigma^2}(y_i - \beta_0 - \beta_1 x_i)^2\right)$$

where  $y_i$  is the log(count) for observation  $i$  and  $x_i$  is the log(weight) for observation  $i$ . Note that here we think of  $\mathbf{y}$  as being random and  $\mathbf{x}$  as being fixed. We could, in theory, think of the vector  $\mathbf{x}$  as also being random and put a prior on it. But we won't do that here.

Now, let's just put uniform priors on  $\beta_0$  and  $\beta_1$  so the priors are proportional to 1. Also, let's assume  $\sigma^2 = 1$ . This seems reasonable since  $s_e^2$ , the MSE, is 0.877. Of course, we could put a prior on  $\sigma^2$  as well and sample it too, but we will focus on only sampling  $\beta_0$  and  $\beta_1$ .

Now, with those uniform priors, and plugging in 1 for  $\sigma^2$ , we have that the joint posterior of  $\beta_0$  and  $\beta_1$  is:

$$p(\beta_0, \beta_1|\mathbf{x}, \mathbf{y}) \propto \exp\left(-\frac{1}{2} \sum_{i=1}^n (y_i - \beta_0 - \beta_1 x_i)^2\right) = f(\beta_0, \beta_1|\mathbf{x}, \mathbf{y}).$$

Then, we can take the log to get

$$\ln(f(\beta_0, \beta_1|\mathbf{x}, \mathbf{y})) = -\frac{1}{2} \sum_{i=1}^n (y_i - \beta_0 - \beta_1 x_i)^2.$$

Our goal now is to obtain samples of  $\beta_0$  and  $\beta_1$ . Let's use the Metropolis algorithm to do this. Using the log of the function proportional to the joint posterior of  $\beta_0$  and  $\beta_1$ ,  $\ln(f(\beta_0, \beta_1|\mathbf{x}, \mathbf{y}))$ , write a Metropolis algorithm in **R**. For  $\beta_0$ , you can use a normal proposal distribution centered at the previous value,  $\beta_0^{(i)}$ , with a standard deviation of 0.8 and for  $\beta_1$ , you can use a normal proposal distribution centered at the previous value,  $\beta_1^{(i)}$ , with a standard deviation of 0.1. The starting values don't matter too much, but we can use  $\beta_0^{(0)} = 10$  and  $\beta_1^{(0)} = -0.5$ . It may be useful to look at the **Notes 8 Script.R** file that is on GitHub in the Notes 8 folder and is on Canvas.

Obtain at least 10,000 samples (set a seed, please) and plot the chains for  $\beta_0$  and  $\beta_1$ . For this problem, include:

1. The plot for the  $\beta_0$  chain.
2. The plot for the  $\beta_1$  chain.
3. The 95% credible interval for  $\beta_0$  based on the 2.5th and 97.5th percentiles.
4. The 95% credible interval for  $\beta_1$  based on the 2.5th and 97.5th percentiles.

```
set.seed(1)
nsamps <- 10000

beta1 <- rep(0, nsamps)
beta0 <- rep(0, nsamps)

beta1[1] <- -0.5
beta0[1] <- 10
```

```

n <- length(tree)
x <- tree$weight
y <- tree$count


for(i in 1:(nsamps - 1)){
  star0 <- rnorm(1, beta0[i], 0.8)
  star1 <- rnorm(1, beta1[i], 0.1)

  logF1 <- -1/2*sum((y - star0 - star1*x)^2)
  logF2 <- -1/2*sum((y - beta0[i] - beta1[i]*x)^2)


  if (log(runif(1)) < (logF1 - logF2)){
    beta0[i + 1] <- star0
    beta1[i + 1] <- star1
  } else{
    beta0[i + 1] <- beta0[i]
    beta1[i + 1] <- beta1[i]
  }
}

```

```
plot(beta0, type = 'l')
```



```
plot(beta1, type = 'l')
```



```
int0 <- quantile(beta0, c(.025, .975))
int0
```

```
##      2.5%      97.5%
## 8.545065 10.211279
```

```
int1 <- quantile(beta1, c(.025, .975))
int1
```

```
##      2.5%      97.5%
## -0.6629883 -0.3576381
```

### Part c (3 points)

Based on the plots of the chains from part b, does it look like the Metropolis sampling worked fairly well?

**They do look pretty random, which is good. They're not getting stuck at any points. I feel like it could be noisier but I think they're perfectly acceptable.**

### Part d (4 points)

Interpret both of the credible intervals from part b.

**There is a 95% chance that the true  $\beta_0$  is between 8.545 and 10.211.**

**There is 95% chance that the true  $\beta_1$  is between -0.663 and -0.358**



### Part e (5 points)

Find and report the integrated autocorrelation time for the  $\beta_0$  and  $\beta_1$  chains. Each chain will have their own  $\hat{\tau}_{int}$  value, so you should report two (although they will be similar).

```
tInt0 <- 1 + 2 * sum(abs(acf(beta0, lag.max = 100, plot = F)$acf))
tInt1 <- 1 + 2 * sum(abs(acf(beta1, lag.max = 100, plot = F)$acf))
tInt0
```

```
## [1] 22.44775
```

```
tInt1
```

```
## [1] 28.65064
```

### Part f (3 points)

Based on the integrated autocorrelation time for the  $\beta_0$  and  $\beta_1$  chains, how many MCMC samples would you need to generate to get the equivalent of 10,000 independent samples?

```
print(tInt0 * 10000)
```

```
## [1] 224477.5
```

```
print(tInt1 * 10000)
```

```
## [1] 286506.4
```

To get the equivalent of 10,000 independent samples for  $\beta_0$  we would need 224478 samples and for  $\beta_1$  we would need 286507 samples.

### Part g (3 points)

Let's compare these credible intervals to some other intervals. First, obtain the 95%  $t$  confidence intervals for  $\beta_0$  and  $\beta_1$  just using the `confint()` function and report them here.

```
tconf <- confint(treeMod, level = 0.95)
tconf0 <- tconf["(Intercept)", ]
tconf1 <- tconf["weight", ]
```

### Part h (10 points)

Now let's obtain confidence intervals using bootstrapping in a similar way we did with regularization in Notes 7 and HW 4 (this is known as bootstrapping the cases). Set a seed and then using at least 10,000 bootstrap samples, report the 95% percentile confidence intervals for  $\beta_0$  and  $\beta_1$  using the `quantile()` function on the values of  $\beta_0$  and  $\beta_1$  that you obtained in the bootstrap.

```

set.seed(1)
n <- 10000
betas <- matrix(rep(0, 2*n), nrow = n)

for(i in 1:n){
  index <- sample(1:nrow(tree), nrow(tree), replace = T)
  betas[i, ] <- coef(
    lm(count ~ weight, data = tree[index, ])
  )
}

```

```

bStrap0 <- quantile(betas[,1], c(0.025, 0.975))
bStrap1 <- quantile(betas[,2], c(0.025, 0.975))

```

### Part i (5 points)

Write a couple sentences comparing all of the intervals in parts b, g, and h.

```

ciDf <- data.frame(
  Variable = c("Beta0 (credible)", "Beta1 (credible)",
               "Beta0 (T Interval)", "Beta1 (T Interval)",
               "Beta0 (bootstrap)", "Beta1 (bootstrap)"),
  Lower = c(int0[1], int1[1], tconf0[1], tconf1[1], bStrap0[1], bStrap1[1]),
  Upper = c(int0[2], int1[2], tconf0[2], tconf1[2], bStrap0[2], bStrap1[2])
)

# Print the table
print(ciDf)

```

```

##           Variable      Lower      Upper
## 1  Beta0 (credible)  8.5450650 10.2112793
## 2  Beta1 (credible) -0.6629883 -0.3576381
## 3 Beta0 (T Interval)  8.5410148 10.2170375
## 4 Beta1 (T Interval) -0.6665051 -0.3633051
## 5 Beta0 (bootstrap)  8.3897056 10.3120985
## 6 Beta1 (bootstrap) -0.6682477 -0.3524438

```

The intervals are actually pretty similar, especially the prediction intervals and the T-intervals. They differ only at roughly 3 decimal places, except for the  $\beta_1$  intervals. Even then, they are not too different. The bootstrap intervals are also pretty similar to the T-intervals and the prediction intervals, although not as similar. They seem to agree on where the true values of  $\beta_0$  and  $\beta_1$  are.